

SRFI 67: Compare procedures

SEBASTIAN EGNER AND JENS AXEL SØGAARD
sebastian.egner-AT-philips.com and jensaxel-AT-soegaard.net

August 18, 2005

Contents

| | | |
|-----|----------------------------------------------|----|
| 1 | Abstract and Rationale | 1 |
| 2 | Introduction | 1 |
| 3 | Terminology and Conventions | 2 |
| 4 | Specification | 2 |
| 4.1 | Comparing atoms | 2 |
| 4.2 | Comparing lists and vectors | 3 |
| 4.3 | Comparing pairs and improper lists | 4 |
| 4.4 | The default compare procedure | 4 |
| 4.5 | Constructing compare procedures | 4 |
| 4.6 | Using compare procedures | 5 |
| 5 | The theory of compare functions | 7 |
| 6 | Design Rationale | 9 |
| 7 | Related work | 14 |
| 8 | Reference implementation | 15 |
| | References | 15 |
| | Alphabetic Index | 17 |

1. Abstract and Rationale

This SRFI can be seen as an extension of the standard procedures `=`, `<`, `char<?` etc. of R⁵RS—or even as a replacement. The primary design aspect in this SRFI is the separation of *representing* a total order and *using it*. For representing the order, we have chosen for truly 3-way comparisons. For using it we provide an extensive set of operations, each of which accepts a procedure used for comparison. Since these compare procedures are often optional, comparing built-in types is as convenient as R⁵RS, sometimes more convenient: For example, testing if the integer index i lies in the integer range $\{0, \dots, n-1\}$ can be written as `(<= /<? 0 i n)`, implicitly invoking `default-compare`.

As soon as new total orders are required, the infrastructure provided by this SRFI is far more convenient and often even more efficient than building each total order from scratch.

Moreover, in case Scheme users and implementors find this mechanism useful and adopt it, the benefit of having a uniform interface to total orders to be used in data structures will manifest itself. Most concretely, a new sorting procedure in the spirit of this SRFI would have the interface `(my-sort [compare] xs)`, using `default-compare` if the

optional *compare* was not provided. Then `my-sort` could be defined using the entire infrastructure of this SRFI: Efficient 2- and 3-way branching, testing for chains and pairwise inequality, min/max, and general order statistics.

2. Introduction

This SRFI defines a mechanism for comparing Scheme values with respect to a total order (aka linear order) [1]. The mechanism provides operations for:

1. comparing objects of the built-in types,
2. using a total order in situations that arise in programs,
3. facilitating the definition of a new total order.

In the following, these aspects will briefly be illustrated.

Traditionally, a total order is represented in Scheme by an order predicate, like `<` or `char<?`. For the purpose of this SRFI, however, a total order is represented by a Scheme-procedure comparing its two arguments and returning either `-1`, `0`, or `1` depending on whether the first argument is considered smaller, equal, or greater than the second argument respectively. Examples of such compare procedures include `(lambda (x y) (sign (- x y)))` for comparing real numbers, but also `(lambda (x y) 0)` comparing anything. For most built-in types specified in the Revised⁵ Report on the Algorithmic Language Scheme (R⁵RS, [3]) compare procedures are specified in Sections 4.1, 4.2, and 4.3 of this SRFI. An axiomatic definition of “compare procedure” is given in Section 5.

The primary reason for using 3-valued compare procedures instead of (2-valued) order predicates is efficiency: When comparison is computationally expensive, it is wasteful if *two* predicates are evaluated where a single 3-valued comparison would suffice. This point is discussed in greater detail in Section 6.

But dealing directly with 3-valued comparisons in the application program is inconvenient and obscures intention: For testing `x < y` one would have to write `(eqv? (compare x y) -1)`. For this reason, an operation `<?` is supplied which allows to phrase the same test as `(<? compare x y)`. This is an example of mapping the three possible outcomes of a comparison into the two boolean values

`{#f, #t}`. Since `<?` takes the total order as an explicit parameter, a comfortably large arsenal of tests can be made available for each and every total order (Section 4.6). This deviates from the approach of `R5RS`, in which there are only five operations (`=`, `<`, `>`, `≤`, `≥`)—and for each total order (`real/number`, `char`, `char-ci`, `string`, `string-ci`) a complete set of these five operation is provided.

But still, using `<?` would be inconvenient if the compare procedure would have to be supplied explicitly every time. For this reason, the parameter `compare` is often made optional in this SRFI—and the procedure `default-compare` is used whenever no compare procedure is passed explicitly. `Default-compare` (Section 4.4) defines *some* resonable total order on the built-in types of `R5RS`.

For the third aspect of this SRFI, defining compare procedures, special control structures (macros) are provided (Section 4.5). These control structures can be used in the definition of a (potentially recursive) compare procedure. This is best explained by an extended example.

Example Assume there is a type `length` representing physical length. The type has an accessor procedure `meters` returning the length in meters (a real number). A compare procedure for lengths can then be defined in terms of `real-compare` (Section 4.1) as:

```
(define (length-compare length1 length2)
  (real-compare (meters length1) (meters length2)))
```

Now, `(<? length-compare x y)` tests if length `x` is shorter than length `y`. Also, `(<=/? length-compare a x b)` tests if length `x` lies between length `a` (incl.) and length `b` (excl.) The expression `(min-compare length-compare x y z)` is a shortest of the lengths `x`, `y`, and `z`. Likewise, `(chain<? length-compare x1 x2 x3 x4)` test if the lengths `x1` `x2` `x3` `x3` are strictly increasing, and so on (refer to Section 4.6).

Furthermore, assume there is another type `box` representing a physical box. The type has procedures `width`, `height`, and `depth` accessing the dimension (each giving a `length`). A compare procedure for boxes, comparing first by width then by height and then by depth, can be defined using the control structure `refine-compare` (Section 4.5) as:

```
(define (box-compare box1 box2)
  (refine-compare
    (length-compare (width box1) (width box2))
    (length-compare (height box1) (height box2))
    (length-compare (depth box1) (depth box2))))
```

This time, `(<? box-compare b1 b2)` tests if box `b1` is smaller than box `b2`—in the sense of the order defined. Of course, all the other tests, minimum, maximum etc. are available, too.

As a final complication, assume that there is also a type `bowl` with accessors `radius` (a `length`) and `open?` (a boolean). Bowls are to be compared first by whether they are open or closed, and then by radius. However, bowls and boxes also need to be compared to each other, ordered such that a bowl is considered “smaller” than a box. (There are type-test predicates `box?` and `bowl?`). Using the control structure `select-compare` (Section 4.5) this can be expressed as:

```
(define (container-compare c1 c2)
  (select-compare c1 c2
    (bowl? (boolean-compare (open? c1) (open? c2)))
    (length-compare (radius c1) (radius c2)))
  (box? (box-compare c1 c2))
  (else "neither bowls nor boxes" c1 c2)))
```

This is an example of “hierarchical extension” of compare procedures, as explained in Section 5. Also note the implicit use of `refine-compare` in the `bowl?`-case.

The preceeding example illustrates the main functionality of this SRFI. For other examples, refer to Section 4.4, and to the file `examples.scm` included in the reference implementation.

3. Terminology and Conventions

A *compare procedure* is a Scheme-procedure of two arguments returning an exact integer in `{−1, 0, 1}` such that the valid input values are ordered according to some total order. A compare procedure, together with a set of Scheme values to which it is applicable, represents a compare function as defined in Section 5.

A *comparison* is either an expression applying a compare procedure to two values, or the result of such an expression.

Each operation (macro or procedure) processing the value of a comparison checks if the value is indeed an exact integer in the set `{−1, 0, 1}`. If this is not the case, an error is signalled.

Compare procedures expecting certain types of argument should raise an error in case the arguments are not of this type. For most compare procedures specified in this SRFI, this behavior is required. A compare procedure *compare* can be used for type-checking value *x* by evaluating `(compare x x)`, in case that is desired. This is useful in procedures like `chain<?` which guarantee to check each argument unconditionally.

4. Specification

4.1. Comparing atoms

In this section, compare procedures for most of the atomic types of `R5RS` are defined: Booleans, characters, strings, symbols, and numbers.

As a general naming convention, a procedure named

type-compare-order

compares two object of the type *type* with respect to a total order for which *order* is a mnemonic hint (e.g. `-ci` for case-insensitive). Of course, `-order` may be absent if there is just one order or the order is obvious. It is an error if a compare procedure accepting objects of a certain type is called with one or two arguments not of that type.

`(boolean-compare bool1 bool2)` procedure

Compares two booleans, ordered by `#f < #t`. *Note:* A non-`#f` value is *not* interpreted as a “true value,” but rather an error will be signalled.

`(char-compare char1 char2)` procedure

`(char-compare-ci char1 char2)` procedure

Compare characters as `char<=?` and `char-ci<=?` respectively. The suffix `-ci` means “case insensitive.”

`(string-compare string1 string2)` procedure

`(string-compare-ci string1 string2)` procedure

Compare strings as `string<=?` and `string-ci<=?`. The suffix `-ci` means “case insensitive.” *Note:* `Compare-string` could be defined as

```
(define (string-compare string1 string2)
  (vector-compare-as-list char-compare
    string1 string2
    string-length string-ref))
```

`(symbol-compare symbol1 symbol2)` procedure

Compares symbols as `string<=?` on the names returned by `symbol->string`.

`(integer-compare x y)` procedure

`(rational-compare x y)` procedure

`(real-compare x y)` procedure

`(complex-compare x y)` procedure

`(number-compare x y)` procedure

Compare two numbers. It is an error if an argument is not of the type specified by the name of the procedure.

Complex numbers are ordered lexicographically on pairs (re, im) . For objects representing real numbers $\text{sign}(x - y)$ is computed. The ordering for values satisfying `real?` or `complex?` but not representing a real or complex number should be consistent with procedures `=` and `<` of `R5RS`, and apart from that it is unspecified.

Numerical compare procedures are compatible with the `R5RS` numerical tower in the following sense: If *S* is a subtype of the numerical type *T* and *x*, *y* can be represented both in *S* and in *T*, then `compare-S` and `compare-T` compute the same result. *Note:* Floating point formats usually

include several symbolic values not simply representing rational numbers. For example, the IEEE 754 standard defines `-0`, `-Inf`, `+Inf`, and `NaN` (“not a number”) for continuing a calculation in the presence of error conditions. The behavior of the numerical comparison operation is unspecified in case an argument is one of the special symbols. *Warning:* The propagation of inexactness can lead to surprises. In a Scheme system propagating inexactness in complex numbers (such as `PLT`, version 208):

```
(complex-compare (make-rectangular (/ 1 3) 1.)
  (make-rectangular (/ 1 3) -1))
⇒ -1
```

At first glance, one might expect the first complex number to be larger, because the numbers are equal on their real parts and the first imaginary part (`1.`) is larger than the second (`-1`). Closer inspection reveals that the decimal dot causes the first real part to be made inexact upon construction of the complex number, and since `(exact->inexact (/ 1 3))` is less than `(/ 1 3)` in the underlying floating point format used, the real parts decide the comparison of the complex numbers.

4.2. Comparing lists and vectors

In this section compare procedures are defined for Scheme lists and vectors—and for objects that can be accessed like lists or like vectors.

An object *x* can be *accessed like a vector* if there are procedures `size` and `ref` such that `(size x)` is a non-negative integer *n* indicating the number of elements, and `(ref x i)` is the *i*-th element of *x* for $i \in \{0, \dots, n - 1\}$. The default vector access procedures are `vector-length` and `vector-ref`.

An object *x* can be *accessed like a (proper) list* if there are procedures `empty?`, `head`, and `tail` such that `(empty? x)` is a boolean indicating that there are no elements in *x*, `(head x)` is the first element of *x*, and `(tail x)` is an object representing the residual elements of *x*. The default list access procedures are `null?`, `car`, and `cdr`.

Independent of the way the elements are accessed, the natural ordering of vectors and lists differs: Sequences are *compared as vectors* if shorter sequences are smaller than longer sequences, and sequences of the same size are compared lexicographically. Sequences are *compared as lists* if the empty sequence is smallest, and two non-empty sequences are compared by their first elements, and only if the first elements are equal the residual sequences are compared, recursively.

```
(vector-compare
  [ compare ] x y [ size ref ])      procedure
(vector-compare-as-list
  [ compare ] x y [ size ref ])      procedure
(list-compare
```

```
[ compare ] x y [ empty? head tail ] procedure
(list-compare-as-vector
 [ compare ] x y [ empty? head tail ] procedure
```

Compare two sequences x and y , using *compare* for comparing elements. The result is an exact integer in $\{-1, 0, 1\}$. If *compare* is not supplied, *default-compare* is used.

The procedure named *access-compare-as-order* accesses the objects like *access* and compares them with respect to the order given by *order*. The names *type-compare* are abbreviations for *type-compare-as-type*.

Examples:

```
(list-compare      '(2) '(1 2))    => 1
(list-compare-as-vector '(2) '(1 2)) => -1
(vector-compare    '#(2) '#(1 2)) => -1
(vector-compare-as-list '#(2) '#(1 2)) => 1
```

4.3. Comparing pairs and improper lists

In this section, compare procedures for Scheme pairs and (possibly) improper lists are defined.

```
(pair-compare-car compare)      procedure
(pair-compare-cdr compare)      procedure
```

Construct a compare procedure on pairs which only uses the car (only the cdr, respectively), and ignores the other. One could define

```
(define (pair-compare-car compare)
  (lambda (x y) (compare (car x) (car y))))
```

Rationale: *Pair-compare-car* can be used to turn a search data structure (e.g. a heap) into a dictionary: Store (**key** . **value**) pairs and compare them using the compare procedure (*pair-compare-car compare-key*).

```
(pair-compare compare-car compare-cdr pair1 pair2)
(pair-compare [ compare ] obj1 obj2)      procedure
```

Compares two pairs, or (possibly improper) lists.

The 4-ary form compares two pairs *pair₁* *pair₂* by comparing their cars using *compare-car*, and if the cars are equal the cdrs are compared using *compare-cdr*.

The 3-ary form compares two objects by type using the ordering of types

null < *pair* < *neither-null-nor-pair*.

Two objects of type *neither-null-nor-pair* are compared using *compare*. Two pairs are compared by using *compare* on the cars, and if the cars are equal by recursing on the cdrs.

The 2-ary form uses *default-compare* for *compare*.

```
(pair-compare '() 'foo)          => -1
(pair-compare '() '(1 . 2))      => -1
(pair-compare '(1 . 2) 'foo)     => -1
(pair-compare 3 4)                => -1
```

4.4. The default compare procedure

It is convenient to have a compare procedure readily available for comparing most built-in types.

```
(default-compare obj1 obj2)      procedure
compares its arguments by type using the ordering
```

null < *pair* < *boolean* < *char* < *string* < *symbol* <
number < *vector* < *other*

Two objects of the same type *type* are compared as *type-compare* would, if there is such a procedure. The type *null* consists of the empty list '(). The effect of comparing two *other* objects or of comparing cyclic structures (made from lists or vectors) is unspecified. (Implementations are encouraged to add comparisons for other built-in types, e.g. records, regexps, etc.) *Rationale:* *Default-compare* refines *pair-compare* by splitting *neither-null-nor-pair*. *Note:* *Default-compare* could be defined as follows (mind the order of the cases!):

```
(define (default-compare x y)
  (select-compare x y
    (null?      0)
    (pair?      (default-compare (car x) (car y))
                (default-compare (cdr x) (cdr y)))
    (boolean?   (boolean-compare x y))
    (char?      (char-compare x y))
    (string?    (string-compare x y))
    (symbol?    (symbol-compare x y))
    (number?    (number-compare x y))
    (vector?    (vector-compare default-compare x y))
    (else (error "unrecognized types" x y))))
```

4.5. Constructing compare procedures

An important goal of this SRFI is to provide a mechanism for defining new compare procedures as conveniently as possible. The syntactic extensions defined in this section are the primary utilities for doing so.

```
(refine-compare <c1> ...)      syntax
```

Syntax: The $\langle c_i \rangle$ are expressions.

Semantics: The arguments $\langle c_1 \rangle \dots$ are evaluated from left to right until a non-zero value is found (which then is the value) or until there are no more arguments to evaluate (in which case the value is 0). It is allowed that there are no arguments at all.

Note: This macro is the preferred way to define a compare procedure as a refinement (refer to Section 5). Example:

```
(define (compare-rectangle r s)
  (refine-compare
    (compare-length (width r) (width s))
    (compare-length (height r) (height s))))
```

```
(select-compare <x1> <x2> <clause1> ...)      syntax
```

Syntax: Each $\langle \text{clause} \rangle$, with the possible exception of the last, is of the form

```
(<type?> <c1> ...)
```

where $\langle \text{type?} \rangle$ is an expression evaluating to a predicate procedure, and $\langle c_i \rangle$ are expressions evaluating to an exact integer in $\{-1, 0, 1\}$. The last $\langle \text{clause} \rangle$ may be an “else clause”, which has the form

```
(else <c1> ...).
```

Semantics: **Select-compare** is a conditional for defining hierarchical extensions and refinements of compare procedures (refer to Section 5). It compares the values of $\langle x_1 \rangle$ and $\langle x_2 \rangle$ by trying the type tests in order, and applies an implicit **refine-compare** on the consequences upon a match.

In more detail, evaluation proceeds as follows: First $\langle x_1 \rangle$ and $\langle x_2 \rangle$ are evaluated in unspecified order, resulting in values x_1 and x_2 , respectively. Then the clauses are evaluated one by one, from left to right.

For clause $((\text{type?}) \langle c_1 \rangle \dots)$, first $\langle \text{type?} \rangle$ is evaluated resulting in a predicate procedure type? and then the expressions $(\text{type? } x_1)$ and $(\text{type? } x_2)$ are evaluated and interpreted as booleans. If both booleans are true then the overall value is **(refine-compare $\langle c_1 \rangle \dots$)**. If only the first is true the result is -1, if only the second is true the result is 1, and if neither is true the next clause is considered. An **else** clause is treated as if both tests were true. If there are no clauses left, the result is 0.

Select-compare evaluates $\langle x_1 \rangle$ and $\langle x_2 \rangle$ exactly once, even in the absence of any clauses. Moreover, each $\langle \text{type?} \rangle$ is evaluated at most once and the resulting procedure type? is called at most twice.

Note: An example of **select-compare** is the definition of **default-compare** given above.

```
(cond-compare <clause1> ...)      syntax
```

Syntax: Each $\langle \text{clause} \rangle$, with the possible exception of the last, is of the form

```
((<t1> <t2>) <c1> ...)
```

where $\langle t_1 \rangle$ and $\langle t_2 \rangle$ are expressions evaluating to booleans, and $\langle c_i \rangle$ are expressions evaluating to an exact integer in $\{-1, 0, 1\}$. The last $\langle \text{clause} \rangle$ may be an “else clause”, which has the form

```
(else <c1> ...).
```

Semantics: **Cond-compare** is another conditional for defining hierarchical extensions and refinements of compare procedures (refer to Section 5).

Evaluation proceeds as follows: The clauses are evaluated one by one, from left to right. For clause $((\langle t_1 \rangle \langle t_2 \rangle) \langle c_1 \rangle \dots)$, first $\langle t_1 \rangle$ and $\langle t_2 \rangle$ are evaluated and the results are interpreted as boolean values. If both booleans are true then the overall value is **(refine-compare $\langle c_1 \rangle \dots$)**. If only the first is true the result is -1, if only the second is true the result is 1, and if neither is true the next clause is considered. An **else** clause is treated as if both booleans were true. If there are no clauses left (or there are no clauses to begin with), the result is 0.

Cond-compare evaluates each expression at most once.

Rationale: **Cond-compare** and **select-compare** only differ in the way the type tests are specified. Both ways are equivalent, and each way is sometimes more convenient than the other.

4.6. Using compare procedures

The facilities defined in this section provide a mechanism for using a compare procedure (passed as a parameter) in the different situations arising in applications.

```
(if3 <c> <less> <equal> <greater>)      syntax
```

Syntax: $\langle c \rangle$, $\langle \text{less} \rangle$, $\langle \text{equal} \rangle$, and $\langle \text{greater} \rangle$ are expressions.

Semantics: **If3** is the 3-way conditional for comparisons. First $\langle c \rangle$ is evaluated, resulting in value c . The value c must be an exact integer in $\{-1, 0, 1\}$, otherwise an error is signalled. If $c = -1$ then the value of the **if3**-expression is obtained by evaluating $\langle \text{less} \rangle$. If $c = 0$ then $\langle \text{equal} \rangle$ is evaluated. If $c = 1$ then $\langle \text{greater} \rangle$ is evaluated.

Note: As an example, the following procedure inserts x into the sorted list s , possibly replacing the first equivalent element.

```
(define (insert compare x s)
  (if (null? s)
      (list x)
      (if3 (compare x (car s))
            (cons x s)
            (cons x (cdr s)) ; replace
            (cons (car s) (insert compare x (cdr s))))))
```

Rationale: **If3** is the preferred way of branching on the result of a comparison in case all three branches are different.

```
(if=? <c> <consequent> [ <alternate> ])      syntax
(if<? <c> <consequent> [ <alternate> ])      syntax
(if>? <c> <consequent> [ <alternate> ])      syntax
(if<=? <c> <consequent> [ <alternate> ])      syntax
```

```
(if>=? <c> <consequent> [ <alternate> ])    syntax
(if-not=? <c> <consequent> [ <alternate> ])    syntax
```

Syntax: $\langle c \rangle$, $\langle \text{consequent} \rangle$, and $\langle \text{alternate} \rangle$ are expressions. If $\langle \text{alternate} \rangle$ is not provided, `(if #f #f)` is used.

Semantics: These six macros are 2-way conditionals for comparisons. First $\langle c \rangle$ is evaluated, resulting in value c . The value c must be an exact integer in $\{-1, 0, 1\}$, otherwise an error is signalled. Then, depending on the value of c and the name of the macro, either $\langle \text{consequence} \rangle$ or $\langle \text{alternate} \rangle$ is evaluated, and the resulting value is the value of the conditional expression.

The branch is chosen according to the following table:

| | $\langle \text{consequent} \rangle$ | $\langle \text{alternate} \rangle$ |
|-----------------------|-------------------------------------|------------------------------------|
| <code>if=?</code> | $c = 0$ | $c \in \{-1, 1\}$ |
| <code>if<?</code> | $c = -1$ | $c \in \{0, 1\}$ |
| <code>if>?</code> | $c = 1$ | $c \in \{-1, 0\}$ |
| <code>if<=?</code> | $c \in \{-1, 0\}$ | $c = 1$ |
| <code>if>=?</code> | $c \in \{0, 1\}$ | $c = -1$ |
| <code>if-not=?</code> | $c \in \{-1, 1\}$ | $c = 0$ |

Note: The macros `if<=?` etc. are the preferred way of 2-way branching based on the result of a comparison.

```
(=? [ compare ] [ x y ])    procedure
(<? [ compare ] [ x y ])    procedure
(>? [ compare ] [ x y ])    procedure
(<=? [ compare ] [ x y ])    procedure
(>=? [ compare ] [ x y ])    procedure
(not=? [ compare ] [ x y ])    procedure
```

If the values x and y are given, test if x and y are in the relation specified by the name of the procedure $rel?$, with respect to compare procedure $compare$; otherwise construct a predicate procedure.

In the forms $(rel? [compare] x y)$, the result is a boolean (either `#t` or `#f`) depending on $(compare x y)$ and the test $rel?$ as specified for `if<?` etc. If $compare$ is not supplied, `default-compare` is used.

In the form $(rel? [compare])$, the predicate procedure $(\lambda (x y) (rel? compare x y))$ is constructed. Again, if $compare$ is not supplied, `default-compare` is used.

A few examples for illustration

```
(>? "laugh" "LOUD") => #t
(<? string-compare-ci "laugh" "LOUD") => #t
(define char<=? (<=? char-compare))
(sort-by-less '(1 a "b") (<?)) => '("b" a 1)
(sort-by-less '(1 a "b") (>?)) => '(1 a "b")
```

Warning: A common mistake is writing $(<=? x y z)$ where $(<=/? x y z)$ is meant; this will most likely manifest itself at the time the expression $(x y z)$ is evaluated.

```
(<=/? [ compare ] [ x y z ])    procedure
(<=/? [ compare ] [ x y z ])    procedure
```

```
(<=/? [ compare ] [ x y z ])    procedure
(<=/? [ compare ] [ x y z ])    procedure
(>/? [ compare ] [ x y z ])    procedure
(>/? [ compare ] [ x y z ])    procedure
(>=? [ compare ] [ x y z ])    procedure
(>=? [ compare ] [ x y z ])    procedure
```

Test if x , y , and z form a chain with the two relations specified by the name of the procedure $rel1/rel2?$, with respect to the compare procedure $compare$.

If $compare$ is not provided, `default-compare` is used. If $x y z$ are not provided, a predicate procedure of three arguments is constructed. The order in which the values are compared is unspecified, but each value is compared at least once.

Note: $(<=/? \text{real-compare } 0 x 1)$ tests if x is a real number in the half open interval $[0, 1)$.

```
(chain=? compare x1 ...)    procedure
(chain<? compare x1 ...)    procedure
(chain>? compare x1 ...)    procedure
(chain<=? compare x1 ...)    procedure
(chain>=? compare x1 ...)    procedure
```

Test if the values $x_1 \dots$ (zero or more values) form a chain with respect to the relation specified by the name of the procedure, and with respect to the compare procedure $compare$. The result is a boolean (either `#t` or `#f`). The order in which the values are compared is unspecified, but each value is compared at least once (even if there is just one.)

A sequence of values x_1, \dots, x_n forms a chain with respect to the relation $rel?$ if $(rel? compare x_i x_j)$ for all $1 \leq i < j \leq n$. In particular, this is the case for $n \in \{0, 1\}$.

Since the relations $=, <, >, \leq$, and \geq are transitive, it is sufficient to test $(rel? compare x_i x_{i+1})$ for $1 \leq i < n$.

Note: The reason every x_i participates in at least one comparison is type-checking: After testing if the values form a chain, these value may be assumed to be of the type comparable by $compare$ —and this holds irrespectively of the number of values, or whether they form a chain.

```
(pairwise-not=? compare x1 ...)    procedure
```

Tests if the values $x_1 \dots$ (zero or more values) are pairwise unequal with respect to the compare procedure $compare$. The result is a boolean (either `#t` or `#f`). The order in which the values are compared is unspecified, but each value is compared at least once (even if there is just one).

The values x_1, \dots, x_n are pairwise unequal if $(not=? compare x_i x_j)$ for all $i \neq j$. In particular, this is the case for $n \in \{0, 1\}$.

Since $compare$ defines a total ordering on the values, the property can be checked in time $O(n \log n)$, and implementations are required to do this. (For example by first sorting and then comparing adjacent elements).

(**min-compare** *compare* x_1 x_2 ...) procedure
 (**max-compare** *compare* x_1 x_2 ...) procedure

A minimum or maximum of the values x_1 x_2 ... (one or more values) with respect to the *compare* procedure *compare*.

The result is the first value that is minimal (maximal, respectively). The order in which the values are compared is unspecified, but each value is compared at least once (even if there is just one value).

(**kth-largest** *compare* k x_0 x_1 ...) procedure

The k -th largest element of values x_0 x_1 ... (one or more values) with respect to the *compare* procedure *compare*.

More precisely, (**kth-largest** *compare* k x_0 ... x_{n-1}) returns the (modulo k n)-th element of the unique sequence obtained by stably sorting $(x_0 \dots x_{n-1})$. (Recall that a sorting algorithm is *stable* if it does not permute items with equal key, i.e. equivalent w.r.t. *compare*).

The argument k is an exact integer, and $n \geq 1$. The order in which the values x_i are compared is unspecified, but each value is compared at least once (even if there is just one value).

Note: The 0-th largest element is the minimum, the (-1) -st largest element is the maximum. The median is the $(n-1)/2$ -th largest element if n is odd, and the average of the $(n/2 - 1)$ -st and $n/2$ -th largest elements if n is even.

(**compare-by**< *lt-pred* [x y]) procedure
 (**compare-by**> *gt-pred* [x y]) procedure
 (**compare-by**<= *le-pred* [x y]) procedure
 (**compare-by**>= *ge-pred* [x y]) procedure
 (**compare-by**<=> *eq-pred* *lt-pred* [x y]) procedure
 (**compare-by**>=> *eq-pred* *gt-pred* [x y]) procedure

If optional arguments x and y are present then these are compared with respect to the total order defined by the predicate(s) given; the result is in $\{-1, 0, 1\}$. If x and y are not present then a procedure comparing its two arguments using the predicate(s) given is constructed and returned.

The predicate procedures mean the following: (*lt-pred* x y) tests if $x < y$, (*le-pred* tests for \leq , (*gt-pred* for $>$, (*ge-pred* for \geq , and (*eq-pred* tests if x and y are equivalent. The result returned by a predicate procedure is interpreted as a Scheme truth value (i.e. **#f** is false and non-**#f** is true).

The purpose of the procedures **compare-by***predicate*(s) is to define a compare procedure from an order predicate, and possibly an additional equivalence predicate. If an equivalence predicate *eq-pred* is given, it is called *before* the order predicate because the equivalence may be coarser than the total ordering, and it may also be cheaper.

Note: **Char-compare** could be defined in terms of **char**<=? as

(**define char-compare** (compare-by<=? char<=?))

(**debug-compare** *compare*) procedure

Constructs a compare procedure equivalent to *compare* but with debugging code wrapped around the calls to *compare*. The debugging code signals an error if it detects a violation of the axioms of a compare function. For this it is assumed that *compare* has no side-effects.

More specifically, (**debug-compare** *compare*) evaluates to a compare procedure *compare*₁ which checks reflexivity, antisymmetry, and transitivity of *compare* based on the arguments on which *compare*₁ is called:

The procedure *compare*₁ checks reflexivity on any value passed to *compare*, antisymmetry on any pair of values on which *compare* is called, and transitivity on triples where two of the arguments are from the current call to *compare*₁ and the third is a pseudo-random selection from the two arguments of the previous call to *compare*₁.

Rationale: The test coverage is partial and determined pseudo-randomly, but the execution time of *compare*₁ is only a constant factor larger than the execution time of *compare*.

5. The theory of compare functions

This section contains a theoretical justification for the concept “compare function”. First an axiomatic definition of compare functions is given. Then it is proved that compare functions are just an unconventional way of defining total orders on equivalence classes of elements—and mathematically that is all there is to say about compare functions.

At this point, a mathematician may wonder why we introduce compare functions in the first place. The answer is: Because they are convenient and efficient for writing programs involving total orders.

In order to make this SRFI as accessible as possible we give the theorems and proofs explicitly, no matter how trivial they are.

Definition: A *compare function* on a set X is a function $c : X \times X \rightarrow \{-1, 0, 1\}$ such that for all $x, y, z \in X$

- (R) $c(x, x) = 0$,
- (A) $c(x, y) + c(y, x) = 0$,
- (T) $c(x, y) \leq 0$ and $c(y, z) \leq 0$ implies $c(x, z) \leq 0$.

We call the properties (R) *reflexivity*, (A) *antisymmetry*, and (T) *transitivity*. ■

The archetypical compare function is

$$\begin{aligned} c_{\mathbb{R}} : \mathbb{R} \times \mathbb{R} &\rightarrow \{-1, 0, 1\} \\ (x, y) &\mapsto \text{sign}(x - y); \end{aligned}$$

it compares real numbers with respect to their canonical order. Obviously, $x < y$ if and only if $c(x, y) < 0$, which we will fix as our sign convention: Instead of writing “ $c(x, y) < 0$ ” we will often simply write “ $x < y$ ” when the compare function c is obvious from the context. (And of course, the convention extends to $\leq, =, \geq$, and $>$ in the obvious way.)

The first theorem states that each compare function gives rise to an equivalence relation in a natural way.

Theorem: Let c be a compare function on X . Then the relation \sim defined by

$$x \sim y \quad :\Leftrightarrow \quad c(x, y) = 0,$$

for $x, y \in X$, is an equivalence relation on X .

Proof: Recall that an equivalence relation is reflexive, symmetric, and transitive [2]. We check:

“Reflexive”: Consider $x \in X$. By (R) $c(x, x) = 0$, so $x \sim x$.

“Symmetric”: Consider $x, y \in X$ such that $x \sim y$. By definition of \sim we have $c(x, y) = 0$. By (A) this implies $c(y, x) = -c(x, y) = 0$. Thus $y \sim x$.

“Transitive”: Consider $x, y, z \in X$ such that $x \sim y$ and $y \sim z$. This means $c(x, y) = c(y, z) = 0$. By (T) this implies $c(x, z) \leq 0$. Moreover, by symmetry also $c(y, x) = c(z, y) = 0$ and by (T) this implies $c(z, x) \leq 0$. Hence, $c(x, z) = 0$ meaning $x \sim z$. ■

The next theorem states that the equivalence classes defined by a compare function are also naturally ordered.

Theorem: Let c be a compare function on X and let \sim be the equivalence relation of the previous theorem. We write $[x]$ for the equivalence class containing x , i.e. $[x] = \{y \in X \mid c(x, y) = 0\}$. Then the relation \leq defined by

$$[x] \leq [y] \quad :\Leftrightarrow \quad c(x, y) \leq 0,$$

for $x, y \in X$, is a total order on the set $\{[x] \mid x \in X\}$ of all equivalence classes.

Proof: Recall that a total order relation is reflexive, (weakly) antisymmetric, transitive, and all elements are comparable [1]. Again, we check:

“Reflexive”: Consider $x \in X$. By (R) $c(x, x) = 0$, so $[x] \leq [x]$ for all $x \in X$.

“Antisymmetric”: Consider $x, y \in X$ such that $[x] \leq [y]$ and $[y] \leq [x]$. By definition of \leq , this means $c(x, y) \leq 0$ and $c(y, x) = -c(x, y) \leq 0$ by (A). Hence, $c(x, y) = 0$ which means $[x] = [y]$.

“Transitive”: Consider $x, y, z \in X$ such that $[x] \leq [y]$ and $[y] \leq [z]$. By definition of \leq this means $c(x, y) \leq 0$ and

$c(y, z) \leq 0$. By (T) this implies $c(x, z) \leq 0$ which means $[x] \leq [z]$.

“Comparable”: For $x, y \in X$, $c(x, y) \in \{-1, 0, 1\}$ as c is a compare function. Hence, $c(x, y) \leq 0$, meaning $[x] \leq [y]$, or $c(x, y) \geq 0$, meaning $[y] \leq [x]$ by (A). ■

Finally, the last theorem shows the converse of the previous two: There is a unique compare function for each total order on a set of equivalence classes.

Theorem: Let X be a set, \sim an equivalence relation on X , and \leq a total order on the set of equivalence classes with respect to \sim . Then the function $c : X \times X \rightarrow \{-1, 0, 1\}$ defined by

$$c(x, y) = \begin{cases} -1 & \text{if } [x] \leq [y] \text{ and not } x \sim y, \\ 0 & \text{if } x \sim y, \\ 1 & \text{if } [y] \leq [x] \text{ and not } x \sim y, \end{cases}$$

is a compare function on X giving rise to the order \leq and the equivalence relation \sim .

Proof: First note that c is well-defined as a function, because $[x] \leq [y]$ and $[y] \leq [x]$ imply $[x] = [y]$ (i.e. $x \sim y$) by the fact that \leq is (weakly) antisymmetric. We check the axioms of a compare function:

“(R)”: Reflexivity of \sim implies $c(x, x) = 0$ for all $x \in X$.

“(A)”: Consider $x, y \in X$. Then $[x] \leq [y]$ or $[y] \leq [x]$ because x and y are comparable with \leq . If both properties hold then $[x] = [y]$, meaning $x \sim y$, so $c(x, y) = c(y, x) = 0$. Otherwise, either $c(x, y) = -1$ and $c(y, x) = 1$ or the signs are flipped. In either case, $c(x, y) + c(y, x) = 0$.

“(T)”: Consider $x, y, z \in X$ such that $c(x, y), c(y, z) \leq 0$. Then $[x] \leq [y]$ and $[y] \leq [z]$ by definition of c . Since \leq is transitive, this implies $[x] \leq [z]$, meaning $c(x, z) \leq 0$. ■

At this point the mathematics of compare functions is finished. However, it is instructive to explore constructions making new compare functions from old ones.

Sign flip: Let c be a compare function on X . Then $(x, y) \mapsto -c(x, y)$ is also a compare function on X ; it is identical to $(x, y) \mapsto c(y, x)$.

As it happens, there are only two functions f mapping $\{-1, 0, 1\}$ into itself such that $(x, y) \mapsto f(c(x, y))$ is a compare function if c is one: $f(\gamma) = 0$ and $f(\gamma) = -\gamma$.

Argument transformation: Let c be a compare function on X and consider a function $\varphi : U \rightarrow X$. Then

$$(u, v) \mapsto c(\varphi(u), \varphi(v))$$

is a compare function on the set U .

One could be tempted to consider the case $c(\varphi(u), \psi(v))$, $\varphi \neq \psi$. But this only results in a compare function (i.e. (R), (A), (T) hold) if c , φ , and ψ are closely related.

Refinement: Let $c_{\text{coarse}}, c_{\text{fine}}$ be compare functions on the same set X . Then

$$(x, y) \mapsto \begin{cases} c_{\text{coarse}}(x, y) & \text{if } c_{\text{coarse}}(x, y) \neq 0, \\ c_{\text{fine}}(x, y) & \text{otherwise} \end{cases}$$

is a compare function. By induction, this construction can be repeated a finite number of times, e.g. starting at the coarsest of all compare functions: $(x, y) \mapsto 0$.

Hierarchical extension: Let X, Y be disjoint sets and let c_X, c_Y be compare functions on X, Y , respectively. Then

$$(u, v) \mapsto \begin{cases} c_X(u, v) & \text{if } u, v \in X, \\ -1 & \text{if } u \in X \text{ and } v \in Y, \\ 1 & \text{if } u \in Y \text{ and } v \in X, \\ c_Y(u, v) & \text{if } u, v \in Y \end{cases}$$

is a compare function on $X \cup Y$. The function refines “ $X < Y$ ” by c_X on X and c_Y on Y , respectively. This construction can be generalized to an arbitrary family (mind the axiom of choice) of compare functions on disjoint domains.

In Scheme, this SRFI defines macros **refine-compare**, **select-compare**, and **cond-compare** for providing convenient and efficient ways of defining refinement, hierarchical extension, argument transformation, and sign flip.

6. Design Rationale

In this section we present our reasoning behind the design decisions made for this SRFI. We would like to be explicit on this because we believe that design is not about the outcome of decisions but about the alternatives considered. The section is organized as a Q&A list.

Order predicates (2-way) or 3-way comparisons?

It is mathematical tradition to specify a total order in terms of a “less or equal” (\leq) relation. This usually carries over to programming languages in the form of a $<=$ predicate procedure.

However, there are inherently *three* possible relations between two elements x and y with respect to a total order: $x < y$, $x = y$, and $x > y$. (With respect to a partial order there is a fourth: x and y are uncomparable.) This implies that any mechanism based on 2-valued operations (be it \leq , or $(=, <)$, or other) has cases in which *two* expressions must be evaluated in order to determine the relation between two elements.

In practice, this is a problem if a comparison is computationally expensive. Examples of this are implicitly defined

orders in which the order of elements depends on their relative position in some enumeration. (Think of comparing graphs by isomorphism type.) In this case, each order predicate is as expensive as a compare procedure—implying that a proper 3-way branch could be twice as fast as cascaded 2-way branches. Hence, there is a potentially considerable loss in performance, and it is purely due to the interface for comparisons.

The primary disadvantage of bare 3-way comparisons is that they are less convenient, both in use and in their definition. Luckily, this problem can be solved quite satisfactorily using the syntactic (macro) and procedural abstractions of Scheme (refer to Sections 4.5 and 4.6).

How to represent the three cases?

We have considered the following alternatives for representing the three possible results of a comparison:

1. the exact integers -1, 0, and 1 (used in this SRFI),
2. the sign of an exact immediate integer,
3. the sign of any Scheme number satisfying **real?**,
4. three different symbols (e.g. **'<**, **'=**, and **'>**),
5. an enumeration type consisting of three elements, and
6. a built-in type with self-evaluating special constants (e.g. **#<**, **#=**, and **#>**).

The representation acts as an internal interface between programs comparing objects and programs using these comparisons.

The advantage of using only three values is that the representation of each case is uniquely defined. In particular, this enables the use of **case** instead of **if**, and it ensures portability. Portability of numbers is problematic in R⁵RS due to underspecification and inexactness.

The advantage of using a non-unique (numerical) representation is that the result of a computation can sometimes immediately be used in a branch, much like the “non-**#f** means true”-convention. However, with the operations in Section 4.6 this advantage hardly matters. Moreover, the “non-**#f** means true”-convention is a major cause of unexpected program behavior itself.

The advantage of using $\{-1, 0, 1\}$ over using three symbols is that the integers support additional operations, for example they can directly be used in index computations. A particularly useful operation is **(* sign (compare x y))** which inverts the order relation depending on **sign** (either -1 or 1). In addition, the integers are unique—once it is known that comparisons result in integers it is obvious

which integers. A minor consideration is that Scheme systems usually treat small integers as unboxed values, and that integers are self-evaluating literals.

The advantage of using three symbols is that they can be chosen to be more descriptive. For example, it is more instructive to see (`symbol-compare 'foo 'bar`) result in `'greater` than in 1. Unfortunately, there is no obvious choice of name for the three symbols. Among the choices that make sense are `'less`, `'equal`, `'greater`, or `'lt`, `'eq`, `'gt`, or `'<`, `'=`, `'>`. A disadvantage of using symbols for the three cases is that Scheme symbols are ordered, too, and this ordering may differ from the desired order for the three cases.

Some Scheme implementations provide a mechanism for defining enumeration types. For example `define-enumerated-type` of Scheme 48 can be used to define a type `comparison` consisting of three objects, say `lt`, `eq`, `gt`. The enumeration can also (directly) be defined on top of SRFI 9 (Defining Record Types) [10] by defining three new record types, each of which having a single instance. We regard this approach as preferable over three symbols because comparison results have their own type, and a sufficiently advanced compiler could use this information to eliminate redundant type-checks.

One step further in this direction is the following design alternative we have considered: Due to the fundamental nature of the type `comparison` for programming, it would be worthwhile integrating it into the core language of Scheme. This could take the following form: There are three self-evaluating constants, e.g. written `#<`, `#=`, `#>`, and these are the only instances of the type `comparison`. The type supports two operations: `comparison?` and `comparison-compare`. Furthermore, `eq?`, `eqv?`, and `equal?` need to understand the comparison values. In other words, `comparison` is designed after `boolean`. It is unclear, however, which problem this tight integration of comparisons into the language is solving.

Given this situation, we have chosen for $\{-1, 0, 1\}$, while providing facilities for using this conveniently—in particular it is hardly ever necessary to deal with the integers directly.

How to order complex numbers?

Mathematically, no total order of the complex numbers exists which is compatible with the algebraic or topological structure. Nevertheless, it is useful for programming purposes to have *some* total order of complex numbers readily available.

Several total orders on the complex numbers are at least compatible with the natural ordering of real numbers. The least surprising of these is lexicographic on (re, im) .

How to order special floating point symbols?

Floating point formats often do not only represent rational numbers but extend this set by special symbols, for example `+Inf`, `-Inf`, `NaN` (“Not a number”), and `-0`. How should these symbols be ordered with respect to the ordinary numerical values and with respect to each other? (Refer to the discussion archive starting with <http://srfi.schemers.org/srfi-67/mail-archive/msg00010.html>.)

Let us briefly recall the purpose of the special symbols. The general rationale for introducing special symbols into a floating point format is for numerical calculations to continue in the presence of data-dependent errors, while still retaining some meaningful information about the result. The symbols `+Inf` and `-Inf` indicate that the calculation has produced a value exceeding the representable range. The special symbol `-0`, indicates that a calculation has produced a value of unrepresentable small magnitude, but retains the information that the underflow approached zero from the negative side (otherwise it would be `+0`). This sign information is useful in the presence of branch-cuts. Finally, `NaN` indicates that the information about the value has been lost entirely (example: `-Inf + Inf`) `NaN` avoids raising an exception and allows carrying on with other parts of the calculation. It should be noted that several `NaN`s can exist. For example in the IEEE 754 standard many bit patterns represent `NaN` (whatever the interpretation).

As `+Inf` and `-Inf` are designed to represent extremal numbers, their ordering with respect to real numbers is obvious. For signed zeros, the ordering is also obvious. However, the notion of two zeros (or even three: `-0`, `0`, and `+0`) is incompatible with the arithmetic structure of the real numbers. Hence, in most situations all zeros should be treated as equal, even though this can destroy information about results. But the alternative design may also make sense in certain situations where the full information carried in a floating point object is to be retained.

For `NaN` (or even several `NaN`s) the situation is even more ambiguous because there is not even a natural order relation of `NaN` with the other possible floating point values. One design alternative is to raise an error if `NaN` is to participate in a comparison; the reasoning being “if the control flow depends on a `NaN` you are in trouble anyway”. An alternative is to define some order by force; the reasoning being “if an object satisfies `real?` then it can be compared with `real-compare`.” Neither approach is obviously better than the other.

Given this situation, we have decided to leave the effect of using a special floating point value in `real-compare` unspecified, in line with the approach of R⁵RS. This approach might change once Scheme itself is more explicit about floating point representations and numerical computation.

How to define `default-compare`?

The purpose of `default-compare` is providing *some* well-defined way of comparing two arbitrary Scheme values. This can be used in all situations in which the user is unwilling to define a compare procedure explicitly, for example because the actual details of the total order do not really matter.

As an example, consider the task of dealing with sets of sets of integers. In this case, one could simply use sorted lists without repetition for representing lists and `default-compare` already provides a total order.

However, there are limits as to how `default-compare` can be defined. For example, `default-compare` cannot easily be based on a hash code derived from the pointer representing an object due to the close dependency with the garbage collection mechanism. Also, we believe it to be more useful to applications if `default-compare` is based on type and structure.

Unfortunately, this imposes limits on what can be compared using `default-compare` because it is very desirable to have a portable reference implementation. In particular, portable ways of dealing with circular structures are overly costly.

Naturally, the question arises how the types should be ordered. For this question it is useful to understand that `boolean-compare` and `pair-compare` both already define a total order for all values (at least in principle). Hence, `default-compare` could refine one of them, but unfortunately not both at the same time (unless `#f` and `'()` are minimum and maximum of the order, respectively). Since `pair-compare` is more frequently used than `boolean-compare` we base `default-compare` on `pair-compare`. The other portably comparable types are ordered by increasing complexity, which clearly is an arbitrary choice.

What is the “lexicographic order”?

The *lexicographic order* is a general way of defining an ordering for sequences from an ordering of elements:

In the lexicographic order, the empty sequence is the smallest sequence of all, and two non-empty sequences are first compared by their first element and only if these are equal the residual sequences are compared, recursively.

The lexicographic order has its name from its use in a lexicon: For example, *fun* < *funloving* < *jolly*.

What is the “natural order” of lists and vectors?

By “natural order” of an abstract data type we mean a total order that is defined to match the basic operations supported by the data type.

The basic access operations with constant execution time for Scheme lists are `null?`, `car`, and `cdr`. These are exactly the operations needed for comparing two sequences lexicographically.

The constant time access operations for Scheme vectors are `vector-length` (size) and `vector-ref` (ref). Using these operations, the fundamental ordering of vectors is first comparing by size, and only if the sizes are equal, by comparing the elements lexicographically.

Why are vectors not ordered lexicographically?

In this SRFI, lists and strings are ordered lexicographically (‘LEX’) by default, e.g. `"12" < "2"`. The default order of vectors is first by length and then lexicographically (‘LENGTH-LEX’), e.g. `#(2) < #(1 2)`. Alternatively, vectors could be ordered purely lexicographically, too. In the extreme, lists, strings, and vectors could even be ordered lexicographically as sequences without distinguishing the concrete representation, implying `"12" = (#\1 #\2) = #(#\1 #\2)`.

The choice affects `vector-compare`, `default-compare`, and the way orders are interpreted conceptually. Moreover, this SRFI introduces the terminology “ordered as lists” and “ordered as vectors” to refer to the two fundamental ways of lifting an order to sequences (LEX and LENGTH-LEX). The choice also has implications for any other SRFI introducing container data types (e.g. 66 and 74), in case the author wishes to specify default compare procedures compatible with this SRFI.

Summarizing the discussion, there seem to be three major arguments:

1. Conceptually vectors and lists are representations of sequences, and if there is only one ordering for them it should be LEX.
2. LENGTH-LEX is more fundamental and efficient for types supporting a constant-time ‘size’ operation.
3. Conceptually strings are “vectors of characters” and strings are conventionally ordered LEX by default, so vectors should be ordered LEX as well in order to minimize the potential for confusion.

(Please refer to the discussion archive for details, in particular <http://srfi.schemers.org/srfi-67/mail-archive/msg00054.html>.)

We consider 2. the most important due to its mathematical nature, followed by 1. because it simplifies the design. While this controversial, we think that it is preferable to introduce different orders for different data types, and not derive every order from a single one for sequences. Finally,

we consider 3. a weak argument because the default ordering of strings is motivated primarily historically for ordering written words of (small alphabet) natural languages.

Concerning other vector-like data types, such as those introduced by SRFI 66 and 74, we recommend to define a default ordering which appears most natural for the type. These can conveniently be named *type-as-ordering*. In cases where the order is of minor importance, we recommend to be compatible with this SRFI.

Why so few higher-order constructions?

An alternative for the control structures (macros) **refine-compare**, **select-compare**, and **cond-compare** is a set of higher-order procedures for constructing compare procedures.

We have chosen for control structures instead of higher-order procedures for simplicity. This becomes particularly evident when a recursive compare procedure, e.g. **default-compare**, is to be defined. Using **select-compare** it is possible to define **default-compare** simply as a procedure calling itself in some branches (refer to the example in Section 4.4). In the higher-order approach, the procedure under construction must also be able to call itself, with arguments that are application specific. Expressing this with a flexible higher-order procedure is much more indirect.

Why the operations **<?**, **<=?** etc.?

Programs need both 2-way branching and 3-way branching. For 3-way branching, the conditional **if3** is provided.

For 2-way branching, the set $\{-1, 0, 1\}$ of results of a comparison is mapped onto the set $\{\#f, \#t\}$. There are eight functions from a 3-set into a 2-set; all six non-constant functions are provided as **=?**, **<?**, etc.

The five monotonic functions can be generalized to chains of values. In order to make the compare procedure parameter optional in the ordinary comparisons, separate operations (**chain<?**, **chain<=?** etc.) are defined for chains. For the sixth operation (**not=?**) the generalization to pairwise unequality is defined as **pairwise-not=?**. This operation can be implemented efficiently because the compare procedure also defines a total order.

As chains of length three are still frequently tested in programs (think of a range check “ $0 \leq i < n$ ”), and often two different relations are combined, there are special operations for chains of length three (**</<?**, **</<=?**, etc.)

For convenience, the compare procedure argument is made optional as often as possible. Unfortunately, this opens up a possibility for mistake: Writing (**<=? x y z**) where (**<=/<=? x y z**) is meant. Fortunately, the mistake will likely manifest itself at the time (**x y z**) is evaluated.

Why are **<?** etc. procedures, not macros?

The procedures **<?**, **</<?**, **chain<?** etc. could also have been specified as macros. This would have the advantage that they could make full use of “short evaluation”: A chain of comparisons stops as soon as one of the comparisons has failed; all remaining argument expressions and comparisons need not be evaluated. This is potentially more efficient.

The advantage of procedures, on the other hand, is that in Scheme they are “first class citizens,” meaning that they can be passed as arguments and returned from higher-order procedures.

Taking this approach one step further, one can even require the compare procedures to check the types of all arguments, even if the result of the comparison is already known. This is what Section 6.2.5 of R⁵RS calls “transitive” behavior of the predicates **=**, **<**, etc. For example, (**< 0 x y**) first tests if **x** is positive, and only if this is the case (**< x y**) is tested. But even if **x** is not positive it is checked that **y** is indeed a **real**—otherwise an error is raised. In “short evaluation,” on the contrary, if **x** is not positive, **y** can be an arbitrary Scheme value.

Clearly, “transitive” tests have an overhead, namely that they need to execute potentially redundant type checks. Even worse, as types are only known to the compare procedure the only way to check the type of a value is to compare it, maybe with itself (which should result in 0 by definition of a compare procedure).

The advantage of “transitive” comparisons is the automatic insertion of a type assertion. For example, after (**chain<? integer-compare x y z**) has been evaluated, no matter the result, it is known that **x**, **y**, and **z** are integers. We consider this advantage sufficiently important to pay the price.

Why compare-by< etc.?

It is often easier to define an order predicate, and possibly a separate equivalence relation, than it is to define a compare procedure. For this case, **compare<** etc. provide a convenient and robust way of constructing the associated compare procedure.

As has been learned from writing the reference implementation, despite the fact that each of these procedures is just a few lines of trivial code, they miraculously attract bugs.

How do I define a compare function from just an equivalence?

You better don't.

A compare function defines a total order on equivalence classes, and vice versa (refer to Section 5). Hence, a compare procedure `compare` can be used to test equivalence: `(=? compare x y)`.

In reverse, one could be tempted to define a “compare function” `c` from just an equivalence relation \sim as $c(x, y) = 0$ if $x \sim y$ and $c(x, y) = 1$ otherwise. However, `c` is not antisymmetric (unless all objects are equivalent, i.e. $c(x, y) = 0$ for all x, y) and hence it is not a compare function. In fact, there is no way at all of avoiding a total order on the equivalence classes.

This is also reflected in the fact that there are efficient (log-time) search data structures based on a total order, but we know of no efficient (sublinear worst-case) data structures based solely on an equivalence relation. The following program takes time and space $O(h)$, where h is the number of equivalence classes in use:

```
(define (equal->compare equal)
  (let ((reps '()) (length-reps 0))
    (define (index x)
      (let loop ((i (- length-reps 1)) (rs reps))
        (if (null? rs)
            (let ((i length-reps))
              (set! reps (cons x reps))
              (set! length-reps (+ length-reps 1))
              i)
            (if (equal x (car rs))
                i
                (loop (- i 1) (cdr rs))))))
      (lambda (x y)
        (integer-compare (index x) (index y)))))
```

If `equal` is an equivalence predicate (i.e. it is reflexive, symmetric, and transitive) then `(equal->compare equal)` is a compare procedure for the objects comparable by `equal`. The total order defined is unspecified (as it depends on call sequence).

Note that the equivalence predicate `equal` could be defined by using a *union-find data structure*. But keep in mind that the equivalence relation represented by `equal` must not change while `(equal->compare equal)` is in use—so the union-find data structure must be unite classes.

How do I switch from R⁵RS to this SRFI?

As it happens, the specification of this SRFI is fully compatible with the 25 order predicates found in R⁵RS. The easiest way of switching is by defining the R⁵RS operations in terms of this SRFI. Refer to the file <http://srfi.schemers.org/srfi-67/implementation/r5rs-to-srfi.scm> for the corresponding Scheme-code.

Alternatively, each expression involving a reference to an R⁵RS order predicate can be transformed into an equivalent expression using the facilities of this SRFI. Be reminded though that this requires an understanding of the *context* of the expression in question, in particular variable bindings, macro definitions, and the use of `eval`.

However, if the meaning of an expression may be altered, it is often possible to increase type safety or simplicity. Consider for example the following potential replacements of `(and (<= 0 i) (< i n))`:

```
(and (<=? real-compare 0 i) (<? real-compare i n))
(<=/? real-compare 0 i n) ; always compares n
(<=/? integer-compare 0 i n) ; only integer i, n
(<=/? 0 i n) ; uses default-compare
```

Only the first alternative is equivalent to the original expression, but the other alternatives might be useful, too, depending on the goal.

Why be so tight with types?

Most procedures and macros in this SRFI are required to signal an error if an argument is not according to the type specified, in particular comparison values must be exact integer in $\{-1, 0, 1\}$ at all times. Alternatively, we could have specified that procedures and macros accept values as general as makes sense.

We believe that being tight on types at this fundamental level of a language pays off quickly. In particular, this will simplify debugging. Moreover, static analysis of a program will recognize more variables of a known type, which allows for more unboxed values and tighter compiled code. (Clearly, at the time of this writing this is speculative.)

Is there a performance penalty for this SRFI?

Yes and no.

The focus of the reference implementation is correctness and portability; performance will very likely suffer due to the overhead of internal procedure calls and type-checking.

But as the word “SRFI” suggests, this document is a “request for implementation,” meaning we would love to see this SRFI being implemented efficiently by the implementation experts of particular Scheme systems. In practice, this means that most of the operations defined here, if not all, are supported natively.

In this case, there is no performance penalty for using the mechanisms of this SRFI—using this SRFI might even be faster due to explicit 3-way branching and better typing.

Why are there optional leading arguments?

Some operations have an optional first argument. This is in contrast to common practice in Scheme to put optional arguments after mandatory arguments.

The leading optional argument is always the argument *compare*, representing the total order to be used. If it is missing *default-compare* is used.

In the cases where we have chosen to make *compare* optional it is for the sake of brevity, e.g. in `(<? x y)` instead of enforcing `(<? default-compare x y)`. Although an option introduces potential for confusion (e.g. `(<? x y z)` vs. `(</<? x y z)`), we consider it an important feature for interactive use and convenient programming (e.g. in `(do ((i 0 (+ i 1))) ((=? i n)))`).

Given our decision for optional *compare*, the question arises how to pass the option. In the absence of other widely accepted mechanisms for options, we can only vary the length of the argument list. For historical reasons—before *case-lambda* of SRFI 16—optional arguments are passed at the end of the argument list for simplified parsing. On the other hand, `(<? compare x y)` is more consistent with the rest of the SRFI than `(<? x y compare)`.

Unfortunately, any particular choice here is a compromise, and it is also controversial. (Please refer to the discussion archive for details, in particular <http://srfi.schemers.org/srfi-67/mail-archive/msg00051.html>.) We have chosen for notational convenience in the common case (optional *compare*) and for consistency within this SRFI (leading optional argument).

Why chain<? etc. and not a predicate parameter?

This SRFI specifies the five chain predicates *chain=?*, *chain<?*, *chain>?*, *chain<=?*, and *chain>=?*. An alternative is to define a single chain predicate that has the ordering as a parameter. (Refer to the discussion archive starting with <http://srfi.schemers.org/srfi-67/mail-archive/msg00012.html>.)

The reason we have chosen for five chain predicates is that we use compare procedures to represent orders, not predicate procedures. There are five possible order relations predicates for which a chain test makes sense. (The sixth, *not=?*, is not not transitive and hence requires pairwise testing.) The five chain tests are clearly defined and can be implemented efficiently, their main overhead being the call to the compare procedure.

Why not more higher-order procedures?

In this SRFI *min-compare* accepts a compare procedure as a first mandatory argument, applying the minimum operation to the list of all other arguments. An alternative is

to have *min-compare* accept only the compare procedure (possibly optional) and returning a procedure computing the minimum of all its arguments (with respect to the compare procedure.) In a similar fashion other operations can specified as higher-order procedures.

We have avoided higher-order procedures in this SRFI for simplicity and efficiency. As said repeatedly, compare procedures are the main vehicle to transport total orders from the code site define an order to the code site using an order. Moreover, most operations made available through this SRFI appear rather infrequently in programs, so either way there is little to be gained. Finally, dealing with higher-order procedures often involves writing more parentheses and the more simple-minded Scheme systems will create many short-lived closures.

Why do <? etc. have so many options?

The procedures *=?*, *<?* etc. accept an optional compare procedure but also two optional arguments to compare. This could be made simpler by not specifying some of the cases, or by specifying different procedures for the different functions.

The operations *<?* etc. are the primary mechanism for using compare procedures. As such they should be versatile and concise.

Our original design had two mandatory arguments for objects to compare and an optional argument for the compare procedure, i.e. it provides a parametric comparison `(<? compare x y)` of two objects. Amir Livne Bar-On then raised the issue of having better support for a higher-order style of programming, i.e. `((<? compare) x y)`. (Refer to <http://srfi.schemers.org/srfi-67/mail-archive/msg00012.html>.)

However, in Scheme the higher-order style is less convenient than it is in curried programming languages like Haskell or ML. In practice this manifests itself as follows: The most basic and frequent case of comparing atomic objects with respect to the default ordering would read `((<=?) x y)`, which is just two parentheses short of optimal.

Fortunately, Dave Mason proposed a syntax for resolving the apparent alternative parametric test vs. higher order style. (Refer to <http://srfi.schemers.org/srfi-67/mail-archive/msg00014.html>.) By combining both functionalities into a single procedure, the user can choose the style at any moment.

7. Related work

The use of compare procedures is not new; defining control structures (*if3*, *select-compare* etc.) for dealing with

them efficiently, however, seems to be new (at least we have not seen it before).

Total ordering in `R5RS` is represented by typed order predicates, such as `<=`, `char<=?` etc. Although a “less or equal”-predicate is sufficient to define a total order, `R5RS` defines a complete set of compare predicates (that is `=`, `<`, `>`, `≤`, and `≤`) for the sake of convenience and readability. There are 25 procedures related to total orders in `R5RS`. These are named `(=<|>|<=|>=)` and `(char|string)[-ci](=<|>|<=|>=)?`.

The traditional approach in Scheme to equivalence (“Are two values treated as equal?”) is the fixed set of predicates `eq?`, `eqv?`, and `equal?`. Historically, this approach was motivated by the desire to compare only pointers and avoid structural recursion. This SRFI provides the generalization to arbitrary equivalence relations, provided the equivalence classes are totally ordered.

The Ruby programming language [4] provides a method `<=>` which is a compare procedure in the sense of this SRFI. By (re-)defining this method a total order can be defined for the instances of a class, when compared against other objects. All 2-way comparisons are based on `<=>`, but in Ruby essentially every method can be overloaded.

In the Haskell 98 programming language [6] order predicates and compare functions coexist. The type `Ordering` [6, Sect 6.1.8] is an enumeration of the three symbolic constants `LT`, `EQ`, `GT`. The type class `Ord` [6, Sect 6.3.2] asserts the presence of a total order for a type, provided the type class `Eq` [6, Sect 6.3.1] also asserts the presence of an equivalence. Since the default definition of the method `compare` is in terms of the methods `==` and `<=`, and vice versa, it can be chosen easily how to provide the total order without affecting its pattern of use.

The C function `strcmp` [7] of the “string.h”-library acts as a compare procedure in the sense of this SRFI, although it is specified to return an integer of which only the sign matters. Python [5] has a built-in function `cmp` which is a compare procedure in the sense of this SRFI.

In SRFI-32 (Sort libraries) [13] the total orders used for sorting are represented by a “less than” procedure. The discussion archive [13] contains a short discussion thread on the use of 3-value comparisons under the aspect whether they can be used to improve the sorting algorithm itself.

In the `Galore.plt` library of data structures for PLT Scheme, total orders are represented by the signature definition `(define-signature order^ (elm= elm< elm<=))`.

8. Reference implementation

The reference implementation is contained in the file

<http://srfi.schemers.org/srfi-67/implementation/compare.scm>;

it is implemented in `R5RS` (including hygienic macros) together with SRFI-16 (`case-lambda`) [9] SRFI-23 (`error`) [11] SRFI-27 (`random-integer`) [12].

Test code and examples are collected in

<http://srfi.schemers.org/srfi-67/implementation/examples.scm>;

it requires SRFI-42 (`comprehensions`) [14]. The reference implementation and the testing code have been developed and are known to run under PLT/DrScheme 208p1 [15], Scheme 48 1.1 [16], and Chicken 1.70 [17].

Code defining the order predicates of `R5RS` in terms of this SRFI is in the file

<http://srfi.schemers.org/srfi-67/implementation/r5rs-to-srfi.scm>.

REFERENCES

- [1] E. Weisstein: *Totally Ordered Set*, Mathworld at Wolfram Research. mathworld.wolfram.com/TotallyOrderedSet.html
- [2] E. Weisstein: *Equivalence Relation*, Mathworld at Wolfram Research. mathworld.wolfram.com/EquivalenceRelation.html
- [3] R. Kelsey, W. Clinger, J. Rees (eds.): *Revised⁵ Report on the Algorithmic Language Scheme*, Higher-Order and Symbolic Computation, Vol. 11, No. 1, August, 1998. www.schemers.org/Documents/Standards/R5RS/
- [4] Y. Matsumoto: *Programming Ruby. The Pragmatic Programmer’s Guide*. www.ruby-doc.org/docs/ProgrammingRuby/
- [5] G. van Rossum, F. L. Drake, Jr., (ed.): *Python Library Reference*. Release 2.4 of 30 November 2004. Section 2.1 “built-in functions”. Python Software Foundation. <http://docs.python.org/lib/lib.html>
- [6] S. Peyton Jones (ed.): *Haskell 98 Language and Libraries* The Revised Report, December 2002. <http://www.haskell.org/definition/>
- [7] ANSI-C *ISO/IEC 9899:1999*, published 1 December. <http://www.open-std.org/jtc1/sc22/wg14/www/standards>
- [8] J. A. Sogaard: *Data Structures Galore for PLT Scheme*. <http://planet.plt-scheme.org:80/207.1/docs/soegaard/galore.plt/1/1/doc.txt>

- [9] L. T. Hansen: *SRFI 16 Syntax for procedures of variable arity*.
<http://srfi.schemers.org/srfi-16/>
- [10] R. Kelsey: *SRFI 9 Defining record types*.
<http://srfi.schemers.org/srfi-9/>
- [11] S. Houben: *SRFI 23 Error reporting mechanism*.
<http://srfi.schemers.org/srfi-23/>
- [12] S. Egner: *SRFI 27 Sources of random bits*.
<http://srfi.schemers.org/srfi-27/>
- [13] O. Shivers: *SRFI 32 Sort libraries*. Section “Ordering, comparison functions & stability” and mail-archive msg000{23,24,33}.html. SRFI has been withdrawn July 17, 2003.
<http://srfi.schemers.org/srfi-32/>
- [14] S. Egner: *SRFI 42 Eager comprehensions*.
<http://srfi.schemers.org/srfi-42/>
- [15] *PLT Scheme*.
<http://www.plt-scheme.org/>
- [16] R. Kelsey, J. Rees: *Scheme48, version 1.1*.
<http://s48.org/>
- [17] *Chicken, version 1.70*.
www.call-with-current-continuation.org.

Copyright (c) 2005 Sebastian Egner and Jens Axel Søggaard.

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the “Software”), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED “AS IS”, WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.

ALPHABETIC INDEX

The principal entry for each term, procedure, or keyword is listed first, separated from the other entries by a semicolon.

- </<=? 6
- </<? 6
- <=<=? 6
- <=<? 6
- <=? 6
- <? 6
- =? 6
- >/>=? 6
- >/>? 6
- >=>=? 6
- >=>? 6
- >=? 6
- >? 6

- boolean-compare 3

- chain<=? 6
- chain<? 6
- chain=? 6
- chain>=? 6
- chain>? 6
- char-compare 3
- char-compare-ci 3
- compare-by< 7
- compare-by<= 7
- compare-by=< 7
- compare-by=> 7
- compare-by> 7
- compare-by>= 7
- complex-compare 3
- cond-compare 5

- debug-compare 7
- default-compare 4

- if-not=? 6
- if3 5
- if<=? 5
- if<? 5
- if=? 5
- if>=? 6
- if>? 5
- integer-compare 3

- kth-largest 7

- list-compare 3
- list-compare-as-vector 4

- max-compare 7
- min-compare 7
- not=? 6
- number-compare 3

- pair-compare 4
- pair-compare-car 4
- pair-compare-cdr 4
- pairwise-not=? 6

- rational-compare 3
- real-compare 3
- refine-compare 4

- select-compare 5
- string-compare 3
- string-compare-ci 3
- symbol-compare 3

- vector-compare 3
- vector-compare-as-list 3